

Containerized Filesystems : A Performance Investigation of Containers across Filesystems

Jia R. Wu
University of Waterloo

Jichen Zhao
University of Waterloo

Abstract

Your Abstract Text Goes Here. Just a few facts. Whet our appetites.

Container as a Service (CaaS) platforms are being offered to personal and enterprise applications. In this paper, we attempt to profile the performance of three separate filesystems, ext4, reiserfs, and xfs running ontop of Arch Linux.

1 Introduction

Filesystems are the implementation of specific rules and principles governing how data is stored and retrieved on storage medium. At the time of writing, there are multiple various filesystems supported on the Linux operating system including ext4, reiserfs, xfs and many more. These filesystems all contain central concepts such as blocks and inodes, but differ subtly in their implementations (cite the linux documentation project here <http://www.tldp.org/LDP/sag/html/filesystems.html>).

Containers, also known as standardized units of software, are executables built ontop of a kernel. These containers offer a standardized execution environment agnostic to the environment the container is installed on. In this paper, we chose to examine Docker, a popular and free platform for creating and running containers. Docker has been championed by researchers as a platform which allows for reproducible research (cite <https://arxiv.org/pdf/1410.0846.pdf>). We are interested in researching whether or not container based filesystem operations incur a significant performance overhead relative to native filesystem operations.

2 Methods

2.1 Environment

Jason Add Description of Your Computer Here

2.2 Docker

2.3 Benchmarks

Filebench, a popular filesystem and storage benchmark for I/O profiling is used to profile our three filesystems, as it permits custom and scalable workloads. We utilized five workloads which shipped with Filebench: FILESERVER, VARMAIL, VIDEOSERVER, WEBPROXY and WEBSERVER. Each of these benchmarks were called 5 times after a fresh installation of the filesystem. Temporary files were discarded with "rm -rf" after each benchmark prior to running the next benchmark to clear the cache. The runtime of Filebench benchmarks were fixed at 300 seconds with the exception of WEBPROXY. Unless explicitly mentioned, the benchmarks use the predefined settings given by the authors¹. This was to mitigate any ramp-up effects that might occur with our HDD. We had to disable address space layout randomization by setting `randomize_va_space=0` in `/proc/sys/kernel` so that Filebench would execute properly and give consistent results.

2.3.1 FILESERVER

This workload creates a directory tree, and calls a sequence of creates, deletes, appends, reads and writes on various files. The authors of Filebench describe this workload as being similar to SPECsfs. We configured the FILESERVER workload to output files of 1.2GB in size, with one flow for each stat, delete, read, close, append, open, close, write, and create operations.

2.3.2 VARMAIL

In this benchmark, multiple operations of create-append-sync, read-append-sync, read and delete are run in a single directory. The purpose of these operations are to simulate I/O operations on a mail server.

2.3.3 VIDEOSERVER

This workload simulates a videosever by serving video files from one directory and caching videos in a second directory. Videos were configured to be 3.2GB in size.

2.3.4 WEBPROXY

WEBPROXY simulated I/O on a web proxy server. Multiple files were created, written to, closed, and deleted in parallel. This benchmark was configured to run only for 60 seconds as 300 seconds would cause a segfault. The authors of Filebench are aware of this issue on Github.

2.3.5 WEBSERVER

This benchmark is similar to WEBPROXY, however it runs open-read-close operations on files in a directory tree and has an additional step of appending.

5 Acknowledgments

A polite author always includes acknowledgments. Thank everyone, especially those who funded the work.

6 Availability

All relevant scripts and data can be retrieved from the following repository:

https://github.com/JRWu/fall2017_cs854/~myname/SWIG

Notes

¹The predefined settings can be located here: <https://github.com/filebench/filebench/wiki/Predefined-personalities>

3 Results

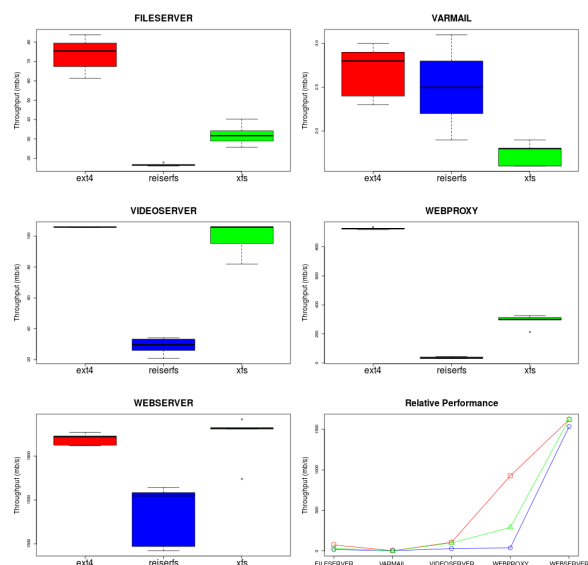


Figure 1: Native FileSystem Performance per Workload

4 Discussion

4.1 Threats To Validity

Well, it's getting boring isn't it. This is the last subsection before we wrap it up.